INVESTMENT DUE DILIGENCE REPORT
ARTIFACT LAB
Confidential | January 2026

INVESTMENT SCORECARD

Rank from A+ to F (with +/-)

Team/Founders – A-

- Veteran gaming infra / monetization team: CEO ex-VP Product & Head of Monetization at King (Activision Blizzard) and VP Data/Targeting at Comcast; VP Product Eng ex-VP AI AdTech at Mistplay; ~20 years combined AAA gaming + adtech experience (Pitch deck – "TEAM" slide)[ppl-ai-file-upload.s3.amazonaws]
- Direct experience building monetization and live-ops systems at King and Mistplay is highly relevant to instrumenting gameplay and extracting behavioral signals at scale (Pitch deck – "TEAM" slide)[ppl-ai-file-upload.s3.amazonaws]
- 2 years of prior R&D on "reasoning data stack" before go-to-market suggests some technical depth but details of IP and team size are not disclosed (Pitch deck – "TEAM" and "Why We'll Win" slides)[ppl-ai-file-upload.s3.amazonaws]
- Key gaps: no visible published ML research profiles or ex-frontier-lab team members; unclear depth in large-scale distributed ML systems vs. data infra and product (Pitch deck – throughout)[ppl-ai-file-upload.s3.amazonaws]

Market Timing – A

- Strong macro tailwind: shift from static LLMs to agentic systems that act and learn in environments; Artifact positions "decision training" as new infra layer spanning pre-training, alignment, evaluation, and post-training loops (Pitch deck – "Decision Training Is A New Infrastructure Layer" slides)[ppl-ai-file-upload.s3.amazonaws]
- Deck frames decision training as a "$10B category" with spend scaling with models × agents × environments × iteration frequency; this is thesis-driven but directionally consistent with rapid funding into AI data & evaluation platforms (Pitch deck – "A $10B Category" and "Recent funding of AI data companies" slides)[ppl-ai-file-upload.s3.amazonaws]
- Gaming behavior data increasingly seen as a high-value asset for training world models and agents (e.g., General Intuition $134M seed at $1B valuation using gameplay data; Scale AI $29B valuation after $1B round; multiple labeling/eval companies valued $1–15B)

(Pitch deck – "Recent funding of AI data companies" slide)[ppl-ai-file-upload.s3.amazonaws]

Problem Clarity – B+

- Core problem: frontier labs and applied-AI teams lack large-scale, structured, human decision trajectories (actions + feedback + outcomes) to train and continuously improve decision-making agents; current text/vision corpora capture content but not "decision loops" (Pitch deck – "AIs next data frontier is here" and "Gameplay is quickly becoming the most prized data source" slides)[ppl-ai-file-upload.s3.amazonaws]
- Clear articulation that gameplay provides dense, observable decision loops (choices, feedback, outcomes) across diverse environments and reward structures vs. sparse or unstructured web data (Pitch deck – "Gameplay matters because it exposes complete decision loops" slide)[ppl-ai-file-upload.s3.amazonaws]
- Less clear who the budget owner is (research vs. product teams vs. infra) and how decision-training spend fits into current lab and enterprise AI spend categories; this is a key GTM risk (Pitch deck – product & business model slides)[ppl-ai-file-upload.s3.amazonaws]

Competition – C+

- Highly capitalized indirect competitors in AI data, labeling, and evaluation: Scale AI, Surge AI, General Intuition, Invisible, Snorkel, Labelbox, Toloka, SuperAnnotate, Mercor; many already provide data, eval, and RL pipelines and can expand into "decision data" (Pitch deck – "Recent funding of AI data companies" slide)[ppl-ai-file-upload.s3.amazonaws]
- Labs also increasingly build internal data-generation and simulation stacks, reducing appetite for vendor lock-in on core training data (industry trend; implied in deck's positioning as infra layer)[ppl-ai-file-upload.s3.amazonaws]
- Artifact claims uniqueness via SDK-level gameplay capture and runtime labeling, but these advantages may be replicable by well-funded incumbents or large publishers if not strongly protected (Pitch deck – "Artifact Difference" and "We turn human play into decisioning data" slides)[ppl-ai-file-upload.s3.amazonaws]

Defensibility – B

- Claimed moat: deep integration with publishers via SDK ("Spark"), consistent and runtime-labeled gameplay events, and equitable rev-share with developers and players, leading to a proprietary "global human decision graph" (Pitch deck – "Technology Moat" and "Artifact Difference" slides)[ppl-ai-file-upload.s3.amazonaws]
- Two-sided network effects: more publishers → richer behavior graph → more valuable to labs and agents; more lab demand → more monetization → more publishers willing to

integrate (Pitch deck – business model and "Human Decision Graph" framing)[ppl-ai-file-upload.s3.amazonaws]

- However, early stage with pipeline but no disclosed long-term exclusive contracts or data rights; multi-homing risk (publishers selling data to multiple buyers) and lab disintermediation risk remain high (Pitch deck – "Pipeline" slides)[ppl-ai-file-upload.s3.amazonaws]

Traction – C

- Supply-side pipeline: 600M+ MAU represented across active publisher pipeline; includes multiple Tier-1 publishers in mobile, PC, cloud gaming with MAUs in 10M–350M range at various stages (alignment, contracting, live) (Pitch deck – "Current momentum – Supply" slide)[ppl-ai-file-upload.s3.amazonaws]
- Demand-side pipeline: one frontier lab actively evaluating data; additional partners across humanoid robotics, embodied AI, and frontier models in "training alignment" or "data evaluation" phases (Pitch deck – "Current momentum – Demand" slide)[ppl-ai-file-upload.s3.amazonaws]
- No disclosed current ARR, contract values, or paid deployments; roadmap targets $10M ARR in 2026 driven by 3–4 foundation labs plus publisher expansion, but current state appears pre-revenue or very early revenue (Pitch deck – "Roadmap – From Pipeline to Scale" and "Reach $10M ARR in 2026" slides)[ppl-ai-file-upload.s3.amazonaws]

Business Model – B

- Two primary revenue engines: (1) model pre-training revenue share tied to usage of decision-grade gameplay signals; (2) post-training SaaS-style recurring revenue for continuous eval, fine-tuning, and decision training of agents and computer-use agents integrated directly into post-training loops (Pitch deck – "Revenue Engines That Compound as We Scale" slides)[ppl-ai-file-upload.s3.amazonaws]
- Pre-training model: rev share (20% cited in deck) from monetization of gameplay signals used in foundation and policy training, scaling with number of publishers and model usage (Pitch deck – "Model Pre-Training – Revenue Share 20%" slide)[ppl-ai-file-upload.s3.amazonaws]
- Post-training model: labs start with paid evaluations and expand into recurring integration for continuous RL/eval; high-margin SaaS potential but highly concentrated, enterprise-style deals (Pitch deck – "Model Post-Training Agents – Recurring SaaS" slide)[ppl-ai-file-upload.s3.amazonaws]

Valuation – N/A (insufficient data; assume high-end seed)

- Deck states a $5M raise "to scale the leading decision-training infrastructure", but does not disclose target valuation; comparable AI data/eval companies are valued from ~$1B to ~$29B, implying that even modest success in this category can support large outcomes (Pitch deck – “Raising $5M…” and “Recent funding of AI data companies” slides)[ppl-ai-file-upload.s3.amazonaws]

- Given pre-revenue / early-revenue status and frontier positioning, likely to be priced at a premium vs. typical infra seed; risk of paying for category story before evidence of repeatable spend (Deck overall)[ppl-ai-file-upload.s3.amazonaws]
- Limited public data on term sheets, past financing, or cap table; need direct confirmation from founders; do not use numbers from unrelated “Artifact AI” accounting automation company (external)[businesswire]

Path to Next Round – B

- Stated target: reach $10M ARR in 2026 via 3–4 foundation lab customers in production plus scaled publisher integrations; this would be strong Series B-level traction if achieved, given strategic nature of accounts (Pitch deck – “Reach $10M ARR in 2026” slide)[ppl-ai-file-upload.s3.amazonaws]
- Plausible Series A path: 1–2 frontier labs in paid production + multiple publishers live with significant MAUs contributing data, and clear unit economics on data/pre-training vs. post-training SaaS (Pitch deck – roadmap slides)[ppl-ai-file-upload.s3.amazonaws]
- High execution risk: long enterprise sales cycles at labs, complex legal/compliance for gameplay data, and need to prove that decision-training data materially improves model performance vs. alternatives; missing this could delay or depress next round (Deck overall)[ppl-ai-file-upload.s3.amazonaws]

---

1. FOUNDERS & TEAM ANALYSIS

[Founder(s) per Deck] – Sony Joseph (CEO); Adam Horwich (VP Product & Engineering)

Background:

- Sony Joseph: former VP Product and Head of Monetization at King (Activision Blizzard) and VP Data & Targeting Models at Comcast; 20 years’ product experience across Ripple, Freewheel, Comcast, and Mint (Pitch deck – “TEAM” slide)[ppl-ai-file-upload.s3.amazonaws]
- Adam Horwich: 15 years’ product experience; former VP AI Ad Tech at Mistplay and 9 years at Activision Blizzard King, where he built King’s monetization platform from 0→1 (Pitch deck – “TEAM” slide)[ppl-ai-file-upload.s3.amazonaws]

- Team has deep expertise in gaming monetization, user telemetry, and data-driven optimization inside large game studios, directly relevant to capturing and productizing gameplay decisions (Pitch deck – "Why We'll Win" slide)[ppl-ai-file-upload.s3.amazonaws]

Strengths:

- Direct domain fit: built exactly the kind of telemetry and monetization pipelines Artifact now wants to leverage for AI; this is rare, specific experience (Pitch deck – "TEAM" and "Why We'll Win" slides)[ppl-ai-file-upload.s3.amazonaws]
- Strong industry network across game studios and AI labs cited as a driver of the current publisher and lab pipeline (2 Tier-1 publishers in alignment/contracting; 1 Tier-1 live; 1 frontier lab in data evaluation) (Pitch deck – "Current momentum – Supply/Demand" slides)[ppl-ai-file-upload.s3.amazonaws]
- Two years of R&D on the "reasoning data stack" pre-product suggests the team has iterated on technical and data-model before scaling GTM (Pitch deck – "Why We'll Win" slide)[ppl-ai-file-upload.s3.amazonaws]

Concerns / Questions:

- No detailed view on broader team composition (ML research, infra engineering, sales) beyond two senior product leaders; unclear whether they can execute on heavy infra + lab sales simultaneously (Pitch deck – "TEAM" slide)[ppl-ai-file-upload.s3.amazonaws]
- No explicit mention of prior startup scaling/exit experience; most background is within large enterprises, which may not directly translate to early-stage company building (Pitch deck – "TEAM" slide)[ppl-ai-file-upload.s3.amazonaws]
- Need clarity on hiring plan for ML science, privacy/legal, and enterprise sales to support both publisher and lab GTM in parallel (Deck overall)[ppl-ai-file-upload.s3.amazonaws]

Team Composition Gaps (Key Risks / Further Diligence)

- Limited visibility into current headcount, org structure, and hiring priorities; risk that team is product-heavy and under-resourced in ML research and enterprise sales (Pitch deck – "TEAM" and roadmap slides)[ppl-ai-file-upload.s3.amazonaws]
- No dedicated compliance or data privacy leadership called out, despite handling potentially sensitive telemetry data and working with global publishers (Deck overall)[ppl-ai-file-upload.s3.amazonaws]
- Need to confirm existence of senior technical lead(s) for large-scale data infra (petabyte-scale pipelines, privacy-preserving analytics, etc.) and lab-facing solution architecture (Further diligence required)[ppl-ai-file-upload.s3.amazonaws]

2. PROBLEM STATEMENT & MARKET

The Thesis

- Artifact aims to become the "default infrastructure for training intelligent decision systems" by converting real human gameplay into decision-grade training signals (decisions, feedback, outcomes) for AI models and agents (Pitch deck – "AIs next data frontier is here" and "Solution" slides)[ppl-ai-file-upload.s3.amazonaws]
- Claims that the next AI frontier beyond text and images is human-like decisioning, which requires dense, structured trajectories of choices and outcomes across many environments; gameplay is positioned as the richest such source (Pitch deck – "Evolution of models training data needs" and "Gameplay is quickly becoming the most prized data source" slides)[ppl-ai-file-upload.s3.amazonaws]
- Artifact provides a "decision training layer" that feeds pre-training, post-training, and continuous evaluation loops with human-grounded decision data; spend scales with number of agents and iteration frequency (Pitch deck – "Decision Training Is A New Infrastructure Layer" slides)[ppl-ai-file-upload.s3.amazonaws]

Is the Problem Real? – YES (but still hypothesis-driven)

- Frontier labs increasingly demand high-quality behavioral data and environments to train agents (e.g., world-model and RL research around games and simulations); Artifact's thesis aligns with visible industry direction (Pitch deck – references to General Intuition, OpenAI interest in videogame data)[ppl-ai-file-upload.s3.amazonaws]
- Gameplay naturally encodes repeated decision loops under varied incentives and constraints (e.g., puzzles, RPGs, multi-agent combat), offering training signals for planning, credit assignment, and multi-agent coordination (Pitch deck – "The future – Decentralized Reasoning Network" and genre breakdown slides)[ppl-ai-file-upload.s3.amazonaws]
- However, evidence that this improves production agent performance vs. synthetic or logged real-world data is still early; most proof points cited in deck are third-party funding rounds rather than independent benchmarks of Artifact's own data (Pitch deck – "Recent funding of AI data companies" slide)[ppl-ai-file-upload.s3.amazonaws]

Market Size Analysis (Quant / Assumptions)

- Internal thesis: decision training is a "$10B category" within AI infra; spend becomes recurring and scales with agents and environments (Pitch deck – "A $10B Category" slide)[ppl-ai-file-upload.s3.amazonaws]
- Comparable AI data and labeling/eval platforms have raised large rounds at high valuations—e.g., Scale AI ($29B valuation after $1B+ round), General Intuition ($1B valuation after $134M seed), Surge AI targeting $15B, Invisible at $2B, Snorkel AI at $1.3B post, etc.—illustrating market appetite for data/eval infra even before full maturation (Pitch deck – "Recent funding of AI data companies" slide)[ppl-ai-file-upload.s3.amazonaws]
- Reasonable assumption: if Artifact can secure a small share (e.g., low-single-digit %) of decision-related training/eval spend at a few frontier labs + major agents platforms, this supports $100M+ ARR potential; however this is highly contingent on labs externalizing this spend rather than building in-house (Assumption based on deck comps and category framing)[ppl-ai-file-upload.s3.amazonaws]

Key Market Concern(s)

- Category definition risk: "decision training layer" is not yet a standard budget line item; could be subsumed into existing data-labeling, eval, or simulation budgets, or built internally by labs (Deck overall)[ppl-ai-file-upload.s3.amazonaws]
- Dependence on a small number of large buyers (frontier labs, big robotics/agent companies) leads to high concentration risk and long, lumpy sales cycles (Pitch deck – demand pipeline slides)[ppl-ai-file-upload.s3.amazonaws]
- Gaming data access is constrained—top publishers have their own AI ambitions and may be reluctant to grant broad rights, especially if labs could use that data for game-like products (Pitch deck – "Equitably sourced" and publisher-integration messaging)[ppl-ai-file-upload.s3.amazonaws]

---

3. COMPETITIVE LANDSCAPE – THE EXISTENTIAL THREAT

Artifact's Claimed Differentiation

- Positions itself as "connective tissue between games and AI models": SDK embedded in games captures structured human behavior, transforms it into training signals, and delivers it to labs as decision intelligence (Pitch deck – "How It Works" slide)[ppl-ai-file-upload.s3.amazonaws]

- Claims to deliver uniquely consistent, runtime-labeled gameplay moments, with time-coded metadata capturing context and behavior; contrasted against unstructured open-internet video and ad-hoc recruitment (Pitch deck – "Artifact Difference" slide)[ppl-ai-file-upload.s3.amazonaws]
- Emphasizes "equitably sourced" data with rev-share and incentives for developers and players, aiming to build a sustainable supply pipeline and avoid extractive data practices (Pitch deck – "Artifact Difference" slide)[ppl-ai-file-upload.s3.amazonaws]

Why This May Not Matter (Risks)

- Existing AI data/eval companies (Scale AI, Surge, General Intuition, etc.) already have strong relationships with labs and could add specialized gameplay pipelines by partnering directly with publishers, narrowing Artifact's window of differentiation (Pitch deck – "Recent funding of AI data companies" slide)[ppl-ai-file-upload.s3.amazonaws]
- Labs may prefer in-house data generation (e.g., simulated environments, synthetic agents, proprietary user logs) over buying third-party gameplay, especially for safety-critical or proprietary domains (Market dynamic inferred; not directly evidenced in deck)[ppl-ai-file-upload.s3.amazonaws]
- Data ownership and rights complexity: if Artifact lacks long-term exclusive rights to gameplay decision data, publishers can multi-home to multiple buyers, compressing margins and limiting moat (Further diligence required; risk inferred from business model)[ppl-ai-file-upload.s3.amazonaws]

---

4. TRACTION & CUSTOMER VALIDATION

What They Claim

- Supply: 600M MAU represented across active publisher pipeline, with several Tier-1 publishers spanning 10M–350M MAU in mobile, PC, and cloud gaming at various deal stages (alignment, contracting, live) (Pitch deck – "Current momentum – Supply" slide)[ppl-ai-file-upload.s3.amazonaws]
- One Tier-1 publisher already live with 10M MAU; indicates at least one integration in production and data flowing, though usage/monetization is not quantified (Pitch deck – "Current momentum – Supply" slide)[ppl-ai-file-upload.s3.amazonaws]
- Demand: 1 frontier lab actively evaluating data; additional partners across humanoid robotics, embodied AI, and AI agents in "data evaluation" or "training alignment" phases;

names withheld due to NDAs (Pitch deck – "Current momentum – Demand" slide)[ppl-ai-file-upload.s3.amazonaws]
- Roadmap: phases from Q4 2025 activation (publishers live, first evaluations) → Q1 2026 monetization (paid evaluations, first paid engagements) → Q2–Q3 2026 scale (multi-buyer demand, post-training platform for computer-use agents) (Pitch deck – "Product Roadmap – Current momentum" slide)[ppl-ai-file-upload.s3.amazonaws]

Red Flags

- No disclosed revenue figures, ARPA, or contract sizes for either publishers or labs; all traction framed as pipeline and roadmap milestones rather than booked ARR (Pitch deck – traction/roadmap slides)[ppl-ai-file-upload.s3.amazonaws]
- Customer logos not shown; all lab and many publisher partners remain under NDA; while understandable at early stage, this makes independent validation difficult (Pitch deck – "Customer names withheld…" note)[ppl-ai-file-upload.s3.amazonaws]
- No empirical evidence in deck that models trained with Artifact's data show meaningfully better performance vs. baseline (e.g., win-rates, sample efficiency, safety improvements), which is critical for labs to commit long-term spend (Pitch deck – product/technology slides)[ppl-ai-file-upload.s3.amazonaws]

What Would Be Convincing

- At least one named frontier or tier-1 lab in public case study with clear before/after performance metrics on relevant tasks (e.g., planning, multi-agent coordination, environment generalization) attributable to Artifact data (Desired diligence outcome)[ppl-ai-file-upload.s3.amazonaws]
- Clear revenue metrics: number of paying labs, ACV per lab, publisher rev-share flows, and overall ARR vs. usage volumes (Desired diligence outcome)[ppl-ai-file-upload.s3.amazonaws]
- Evidence of strong publisher retention and willingness to expand integration to more titles/genres after initial deployment (Desired diligence outcome)[ppl-ai-file-upload.s3.amazonaws]

---

5. BUSINESS MODEL & FINANCIALs

Revenue Model

- **Pre-training:** rev share (20% cited) from labs using decision-grade gameplay signals for foundation and policy pre-training; each additional publisher expands the "global human decision graph", increasing data value and potential usage (Pitch deck – "Model Pre-Training – Revenue share 20%" slide)[ppl-ai-file-upload.s3.amazonaws]
- **Post-training/agents:** recurring SaaS revenue from continuous evaluation, fine-tuning, and decision training for AI models and computer-use agents; labs integrate Artifact directly into post-training iteration loops (Pitch deck – "Model Post-Training Agents – Recurring SaaS revenue" slide)[ppl-ai-file-upload.s3.amazonaws]
- **Publisher economics:** publishers gain incremental monetization with "pure incremental revenue – no cannibalization of IAP/ads; zero integration cost; instantly scalable across live network"; Artifact shares revenue from data usage with publishers (Pitch deck – "Artifact Extension – Optional Incremental Revenue Layer" slide)[ppl-ai-file-upload.s3.amazonaws]

Financial Targets

- Stated ambition: reach $10M ARR in 2026 via lab contracts and publisher expansion; scale to 3–4 foundation labs in production and extend product into post-training pipelines (Pitch deck – "Reach $10M ARR in 2026" slide)[ppl-ai-file-upload.s3.amazonaws]
- Fundraising ask: $5M to scale decision-training infrastructure; prior backers include Raptor Digital, IDG Ventures, Cypher Capital, Blizzard Avalanche Ecosystem Fund (Pitch deck – "Raising $5M…" slides)[ppl-ai-file-upload.s3.amazonaws]
- No explicit burn, runway, cash balance, or historical revenue disclosed; must assume current revenue is minimal or non-material (Deck overall)[ppl-ai-file-upload.s3.amazonaws]

Model Concerns

- Revenue model complexity: tying pre-training rev share to data usage may be difficult to meter and audit, especially when data is combined with other sources inside labs (Pitch deck – business model slides)[ppl-ai-file-upload.s3.amazonaws]
- High dependency on a handful of large labs for both pre-training and post-training revenue; small number of deals drives most ARR, increasing concentration risk (Pitch deck – demand pipeline & roadmap)[ppl-ai-file-upload.s3.amazonaws]
- Publisher rev-share economics not detailed (e.g., % split, minimum guarantees), making it hard to assess long-term margin and bargaining power (Deck overall)[ppl-ai-file-upload.s3.amazonaws]

Unit Economics (Speculative)

- Data infra businesses can exhibit high gross margins once data pipelines and infra are built, but this is dependent on: (a) ability to reuse datasets across many buyers; (b) low incremental cost per additional lab/agent (Assumption based on model type)[ppl-ai-file-upload.s3.amazonaws]
- If Artifact captures 20% rev share from data usage, but must pass a significant portion to publishers (e.g., >50%), net take-rate may be modest, requiring very large usage volumes to justify valuation (Assumption; terms not disclosed)[ppl-ai-file-upload.s3.amazonaws]
- Limited public data on COGS (compute/storage), customer acquisition costs, or payback periods; this is a major diligence gap before underwriting significant scale (Deck overall)[ppl-ai-file-upload.s3.amazonaws]

---

6. DEFENSIBILITY & MOATS

Claimed Moats

- Deep game-engine-level integration via SDK ("Spark") enabling capture and annotation at runtime, aligned incentives with publishers (stickier users, deeper engagement) and a publisher AI workbench for deploying reasoning agents (Pitch deck – "Technology Moat" slide)[ppl-ai-file-upload.s3.amazonaws]
- Consistent, runtime-labeled telemetries yield high-quality, comparable datasets across games; seen as superior to scraped or ad-hoc gameplay clips (Pitch deck – "Artifact Difference" slide)[ppl-ai-file-upload.s3.amazonaws]
- Building a decentralized reasoning network: games → Artifact platform → AGI companies; each integration extends data coverage, compounding value (Pitch deck – "Decentralized Reasoning Network" slides)[ppl-ai-file-upload.s3.amazonaws]

Moat Assessment

- If Artifact secures long-term, preferably exclusive or semi-exclusive rights to high-value publishers at engine level, this can form a meaningful data moat that is costly and slow for competitors to replicate (Pitch deck – publisher pipeline slides)[ppl-ai-file-upload.s3.amazonaws]
- However, current stage appears early with 1 live Tier-1 publisher and others in alignment/contracting; no explicit exclusivity or contract lengths are disclosed (Pitch deck – "Current momentum – Supply" slide)[ppl-ai-file-upload.s3.amazonaws]

- Data moats in AI are increasingly contested (labs training on public web, synthetic data, and first-party logs); Artifact must show that its human decision graph is both unique and necessary for key agent capabilities (Deck overall)[ppl-ai-file-upload.s3.amazonaws]

Category vs. Feature Risk

- Category risk: "decision training" as a distinct infra category may not crystallize; could be absorbed into broader AI data/eval platforms or into labs' internal teams, reducing standalone value capture (Deck overall)[ppl-ai-file-upload.s3.amazonaws]
- Feature risk: incumbents like Scale or General Intuition could add "gameplay decision data" as a product feature, leveraging existing distribution and sales relationships with labs (Pitch deck – "Recent funding…" slide)[ppl-ai-file-upload.s3.amazonaws]
- Artifact's best chance at category-defining status is to become synonymous with gameplay-based decision training—requiring fast land-grab of publishers and a few flagship lab wins (Deck overall)[ppl-ai-file-upload.s3.amazonaws]

---

7. CRITICAL RISK FACTORS

Existential Risks

- Failure to secure sufficient high-quality, high-MAU publisher integrations (and associated rights) would leave Artifact with undifferentiated or low-value data vs. large peers (Pitch deck – publisher pipeline slides)[ppl-ai-file-upload.s3.amazonaws]
- Labs decide to generate or source comparable data in-house or via existing partners, rendering Artifact's offering marginal or unnecessary (Pitch deck – competitive/comps slide)[ppl-ai-file-upload.s3.amazonaws]
- Regulatory or platform shifts (privacy laws, platform policies, game EULAs) restrict the collection or commercialization of detailed gameplay telemetry at scale (Regulatory risk; not detailed in deck)[ppl-ai-file-upload.s3.amazonaws]

Execution Risks

- Dual-sided GTM (publishers + labs) is complex; success depends on orchestrating partnerships across two demanding enterprise customer groups with very different motivations and sales cycles (Pitch deck – "How it works" and pipeline slides)[ppl-ai-file-upload.s3.amazonaws]

- Need to prove causal link between Artifact data and improved model/agent performance quickly; otherwise the company risks being perceived as "nice-to-have data exhaust" vs. critical infra (Deck overall)[ppl-ai-file-upload.s3.amazonaws]
- Technical complexity of building robust, privacy-aware, scalable pipelines and tools (SDKs, ingestion, auto-annotation, RL gym abstractions, eval/verifier tooling) within limited seed-stage resources (Pitch deck – "Enabling capabilities" roadmap slide)[ppl-ai-file-upload.s3.amazonaws]

Strategic Risks

- Brand and naming collisions with other "Artifact AI" companies (e.g., Artifact AI accounting-automation platform backed by a16z speedrun, launched in UK/US), which may confuse customers and investors and complicate SEO/brand (BusinessWire/Yahoo coverage; BusinessWire, Yahoo Finance)finance.yahoo+1
- Potential misalignment between publishers (who may fear cannibalization or IP leakage) and labs (who want rich, unconstrained data), complicating data-rights negotiations (Deck overall)[ppl-ai-file-upload.s3.amazonaws]
- Over-reliance on gaming as a sole data domain; if the market shifts toward other environments (e.g., enterprise apps, robotics logs), Artifact must successfully expand beyond gameplay (Pitch deck – "Future expansion – beyond gameplay data" hints)[ppl-ai-file-upload.s3.amazonaws]

---

8. QUESTIONS FOR FOUNDERS

Competition

- Which AI data/eval companies are most often encountered in lab conversations (e.g., Scale, Surge, General Intuition), and how does Artifact win/lose vs. them today? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]
- Are there any existing or planned strategic partnerships with major AI data players to distribute Artifact's data vs. competing head-on? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]

Traction

- What is current ARR / signed contract value, broken down by (a) publishers and (b) labs? How much of this is recurring vs. one-off evaluations? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]
- Can the team share anonymized case studies showing quantifiable model/agent performance improvements attributable to Artifact data? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]
- What is the current conversion rate from data evaluation → paid production deployments at labs? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]

Business Model

- How is the 20% pre-training rev share calculated, enforced, and audited in practice, especially when data is mixed with other sources inside labs? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]
- What is the economic split between Artifact and publishers for data monetization? Are there minimum guarantees or volume tiers? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]
- For post-training SaaS, what is target ACV per lab, and what are expected sales cycles and required professional services? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]

Team

- Beyond the two highlighted leaders, what is the composition of the core team (ML research, infra, sales, legal/compliance)? What are the next 5 critical hires? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]
- Does the team have dedicated in-house expertise in privacy, data protection, and global compliance (GDPR, COPPA, etc.) as relevant to gameplay telemetry? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]

Defensibility

- Are any publisher agreements exclusive (by genre, geography, or data type)? What is the average contract term and renewal profile? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]
- How difficult would it be for a major publisher or AI data incumbent to replicate Artifact's SDK + runtime labeling approach if they chose to? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]

- How does Artifact plan to extend its data moat beyond games (e.g., to enterprise apps, robotics, or other human-computer interaction data)? (Further diligence question)[ppl-ai-file-upload.s3.amazonaws]

---

9. CONDITIONS FOR INVESTMENT

Bull Case (Path to Success)

- Artifact becomes the de-facto standard for gameplay-based decision training, securing long-term, high-MAU integrations with multiple Tier-1 publishers and a handful (3–4) of frontier labs in production by 2026 (Pitch deck – roadmap and pipeline)[ppl-ai-file-upload.s3.amazonaws]
- Demonstrable uplift in agent performance (planning, robustness, multi-agent behavior) from Artifact data leads labs and agents platforms to treat Artifact as critical infra, supporting large, sticky ACVs and multi-year contracts (Desired outcome)[ppl-ai-file-upload.s3.amazonaws]
- Recurring SaaS revenue from post-training/eval loops grows on top of data rev share, driving ARR beyond $10M into $50–100M+ over time with strong gross margins and modest net revenue retention >120% (Assumption based on successful infra pattern)[ppl-ai-file-upload.s3.amazonaws]

Bear Case (Failure Modes)

- Labs and data incumbents develop equivalent or "good enough" decision-training pipelines internally, limiting Artifact to a niche or "nice-to-have" role with small contracts and low pricing power (Risk as in Sections 3 and 6)[ppl-ai-file-upload.s3.amazonaws]
- Publisher integrations stall due to legal/compliance friction, integration complexity, or misaligned incentives; Artifact fails to build a sufficiently rich and proprietary human decision graph (Risk as in Sections 2 and 7)[ppl-ai-file-upload.s3.amazonaws]
- Category framing ("decision training layer") fails to solidify; buyers treat this as generic data labeling/eval, compressing multiples and eroding the narrative that underpins high-valuation outcomes (Risk as in Market section)[ppl-ai-file-upload.s3.amazonaws]

Required Answers Before Committing

- Concrete, up-to-date metrics on: (a) number of live publishers and MAUs under contract; (b) number and nature of paying lab customers; (c) current ARR and pipeline conversion assumptions (Further diligence)[ppl-ai-file-upload.s3.amazonaws]
- Details on data rights (ownership, exclusivity, duration, re-use rights) in publisher contracts and data-use terms with labs (Further diligence)[ppl-ai-file-upload.s3.amazonaws]
- Evidence (even if anonymized) that Artifact's data materially improves model/agent performance vs. alternatives—and how that improvement maps into willingness-to-pay and contract sizes (Further diligence)[ppl-ai-file-upload.s3.amazonaws]

If Investing (Indicative Stance)

- Consider participating or leading **if**: (1) at least one frontier or tier-1 lab is in or near paid production; (2) there are signed contracts with multiple Tier-1 publishers representing >100M MAU with favorable data rights; and (3) there is a clear 18–24 month runway at proposed valuation (Investment condition)[ppl-ai-file-upload.s3.amazonaws]
- Push for board visibility and structured milestones tied to: publisher integrations, lab conversions, and demonstrated model performance gains, with reserves earmarked for potential pro-rata in strong-signal next round (Investment condition)[ppl-ai-file-upload.s3.amazonaws]
- Negotiate protections around potential brand confusion with other "Artifact AI" entities and confirm IP/brand position to avoid downstream legal or market confusion (Investment condition; external comps: BusinessWire)[businesswire]

Note on Materials Used & Assumptions

- This report is based primarily on the **artifact-deck.pdf** (Artifact Lab 2025 pitch deck) and the provided diligence template; limited or no public data was found specifically on "Artifact Lab" as a company beyond the deck contents (Deck; Diligence-Template.docx)Diligence-Template.docx+1
- A separate class of company named "Artifact AI" (agentic AI accountant for accounting firms) exists and is backed by a16z speedrun; this appears to be unrelated and has been treated only as a naming/brand risk, not as core factual input to this diligence (Artifact AI accounting site; BusinessWire)finance.yahoo+1
- Where data is incomplete (revenue, valuation, precise contract terms), explicit assumptions are labeled as such and should be validated directly with the founders.